Databases II

Laboratory Exercise 2020/21

Name	Surname	AM
Nikiforos – George	Papageorgiou	1059633
Nicholas	Stamopoulos	1057764

I certify that I am the author of this work and that I have expressly and specifically cited or referenced in it all sources from which I have used data, ideas, sentences or words, whether they are accurately conveyed (in the original or translated) or paraphrased. I also certify that this assignment was prepared by me personally specifically for this particular course/seminar/curriculum.

I have been informed that according to the internal regulation of the University of Patras, article 50§6, any attempt to copy or in general to tamper with the examination and educational process by any examinee, in addition to nihilism, constitutes a serious disciplinary offense.

Signature

22 / 09 / 2021

Signature

22 / 09 / 2021

Code files attached

We submit the following code files along with this report

file	Regarding question	Description/
queryX.ipynb	1	**Comment** Contains the code for query X, where X ÿ [1, 10]

Technical characteristics of the operating environment

Technical characteristics of the physical PC used for the work

Feature	Price
CPU model	AMD Ryzen 5 2600 Six-Core Processor
CPU clock speed	3.4GHz
Physical CPU cores	6
Logical CPU cores	12
RAM	16Gb
Secondary Storage Type	SSD

Technical characteristics of the virtual machine (VM) used for the job **Feature**

	Price
CPU cores	6
Execution cap	100%
RAM	9Gb
VM OS	Ubuntu 20.04.2 LTS
VM software	VirtualBox Windows 10
Host OS	

Question 1: Question Answers

Note: The requested results are listed in screenshot format, for better illustration.

Query Give	**Answer**
the number of users who watched the movie “Jumanji”. Give the names of	+--------------+-------------+ |title|total_viewers| +--------------+-------------+ |Jumanji (1995)|22243| +--------------+-------------+
movies that users rated as “boring”. Give users who have rated the movie as	+-----------------------------------+---------+ |title|lower_tag| +-----------------------------------+---------+ |(500) Days of Summer (2009)|boring| |101 Reykjavik (101 Reykjavík) (2000)|boring| |12 Years a Slave (2013)|boring| |1408 (2007)|boring| |1492: Conquest of Paradise (1992)|boring|
"Bollywood" and have rated it with a grade >3.	+------+------+-------------------+ |userId|rating|lower_tag| +------+------+-------------------+ | 10573| 4.0|bollywood| | 19837| 5.0|bollywood| | 23333| 4.0|bollywood| | 25004| 5.0|bollywood| | 31338| 4.5|bollywood|

Find the top 10 movies for each year.	
Give the tags for each movie and the name of the movie for the year 2015.	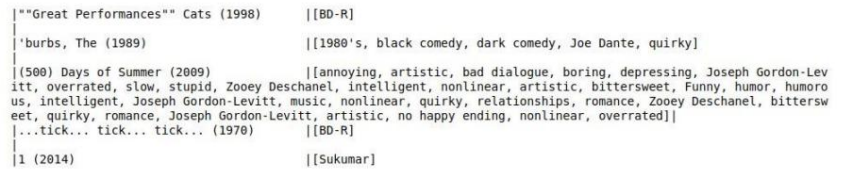
Give the number of ratings for each movie.	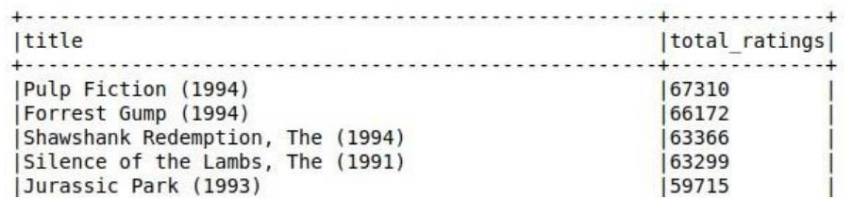
Find the first 10 users with the most ratings for each year. Find	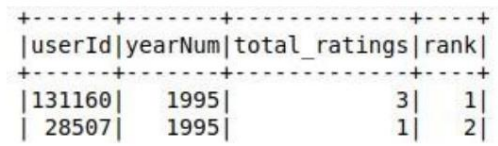
the movies with the most ratings for each movie category.	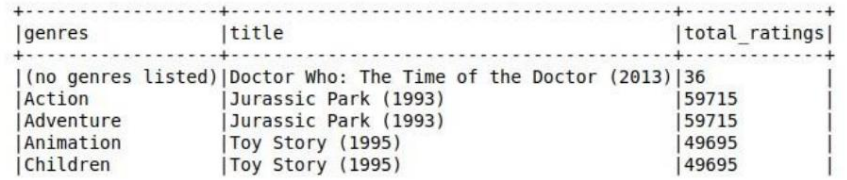
Give the total number of users watching the same movie, on the same day and time. Give the	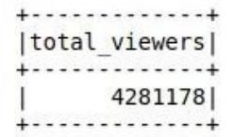
number of movies, for each category, that users rated as "funny" and with a rating > 3.5.	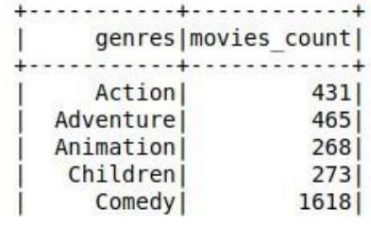

Question 2: Performance comparison on single node/virtual cluster/Livy

Virtual cluster settings

A/A 1	Executor cores Executor mem Driver cores			Driver mem
	1	1G	1	1G
2	2	2G	1	1G
3	2	2G	2	2G

Execution Times

Note: The execution times below were measured in seconds. Timing was done using the sparkMeasure library.

Question 1	Local	Virtual 1	Virtual 2	Virtual 3	Livy
	14	40	28	29	120
2	5	19	15	15	15
3	20	46	40	38	123
4	26	55	44	52	180
5	4	18	15	13	16
6	15	40	36	34	108
7	27	53	41	41	168
8	25	48	41	38	119
9	32	90	62	63	181
10	16	38	30	32	120

Analysis of results

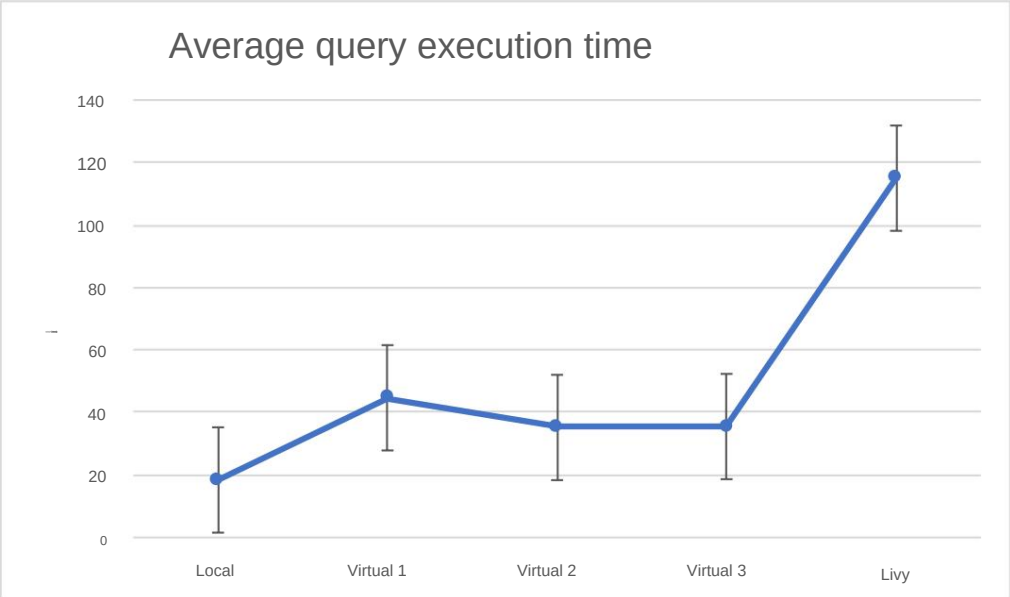

After measuring the time and analyzing the results, we make the following observations:

• On a single node machine (local) we achieve fast query execution, as all the computing power we have assigned to the VM is used. • Between Virtual 1, Virtual 2 and Virtual 3, we notice that the first consumes more time to execute queries than the other two and this is logically due to the assignment of only one core for each worker. We also notice that between Virtual 2 and 3 the differences are negligible, therefore the increase in cores and memory that the driver binds in Virtual 3 did not bring better results than 2. • The execution of queries on the Livy server consumes the maximum time. • Regarding the queries, we notice that specifically queries 2 and 5 always run in much less time compared to the rest. We suspect that this is due to the fact that these two queries do not make use of the rating.csv file, which contains the largest number of records of all and is therefore more "expensive" in terms of operations.

Bibliography

1. *PySpark 3.1.2 Documentation.* http://spark.apache.org/docs/latest/api/python/ 2. A. Komnenos. *Tutorial 6 – Introduction to Apache Spark.* https://eclass.upatras.gr/modules/document/index.php?course=CEID1176&openDir =/5e6f65ear83d/60756472Ab51 3. Nishant Bahri. *Movie Lens Data Analysis Using PySpark [for beginners].* https://medium.com/analytics-vidhya/movie-lens-data-analysis-using-pyspark-for beginners-9c0f5f21eaf5 4. Mauro Krikorian. *Movie Data Statistics with Apache Spark.* https://medium.com/southworks/movie-data-statistics-with-apache-spark 58c2ef8fe452